

Table 1: Local C2P-Cache evaluation setup and default parameters. Values describe the reproducible Accel-Sim configuration, not an assertion of bit-exact equivalence with the paper.

Parameter	Local value
<i>(a) Baseline GPU configuration</i>	
GPU configuration	64 SMs; 1.41 GHz core/L2/interconnect model
Simulated endpoint organization	64 clusters $\times$ 1 SM; logical C2P peer group: 8 SMs
Warp scheduler	GTO; 4 schedulers/SM
Warp size / thread limit	32 threads/warp; 2048 threads/SM
Registers	65,536 registers/SM
Private L1 cache	64 KiB/SM; 16 sets; 32-way; 128 B line; 20-cycle hit latency
Shared L2 cache	128 sets; 16-way; 128 B line; 200-cycle baseline path
Memory system	20 memory partitions; 2 sub-partitions/partition
<i>(b) Modeled C2P timing and bounded resources (cycles)</i>	
Bloom-filter engine / snapshot lookup	2 / 2
Remote tag lookup / base return	7 / 2
Target probe FIFO / abort timeout	32 entries / 32 cycles
Remote request and response paths	Dedicated far-L1 queues; target L1 tag/data port contention modeled
<i>(c) Default C2P configuration</i>	
Bloom-filter size (m)	5,120 bits
Tag-mask / BF-hash region	1,024 / 4,096 bits
Encoding (k)	1 tag-mask + 3 BF-hash positions
Snapshot organization	64 banks; 4 copies/bank (1 original + 3 replicas)
BF engines / update transport	128 engines / 128 B per cycle